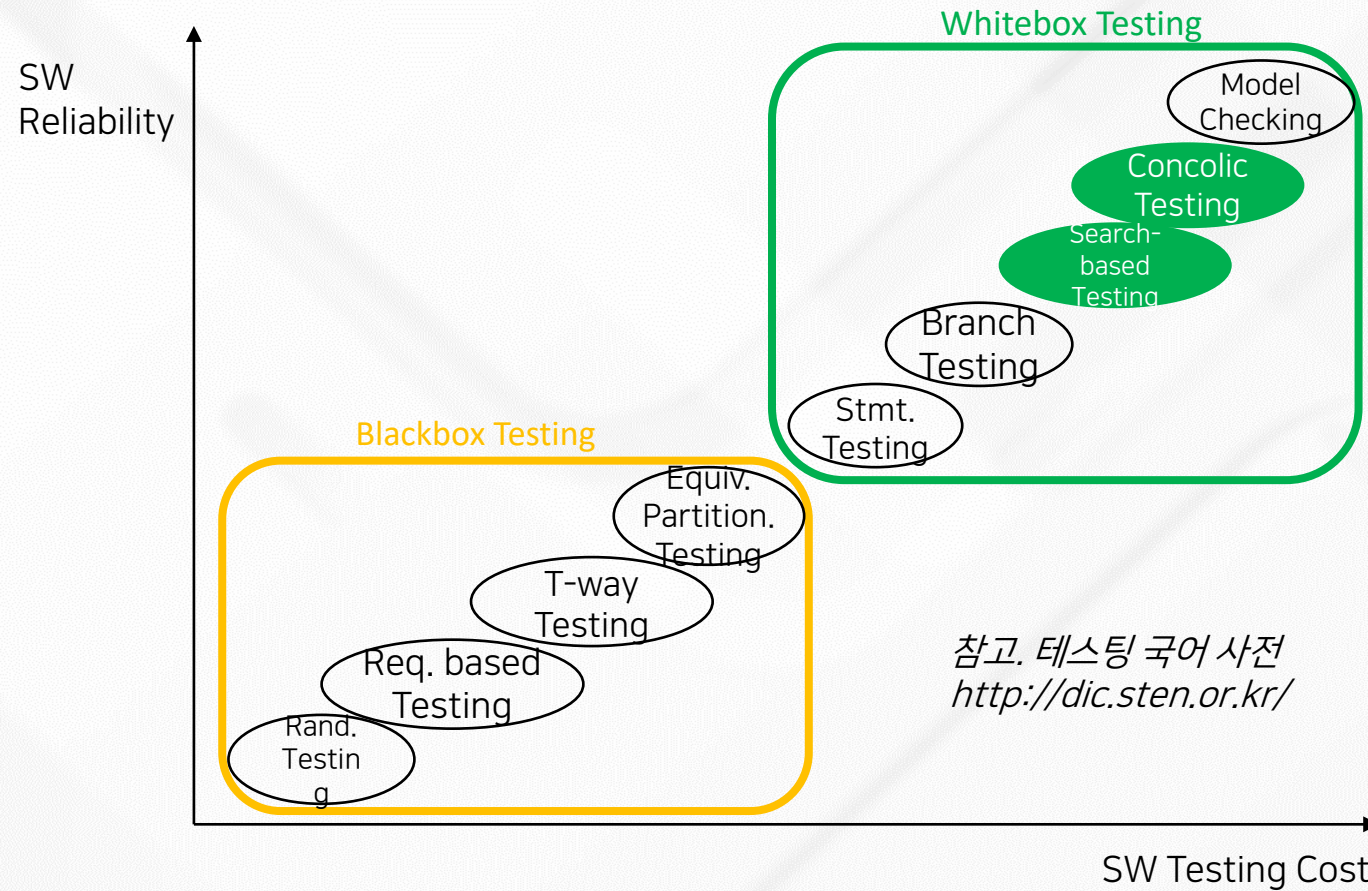


Automated Test Input Generation through Fuzzing

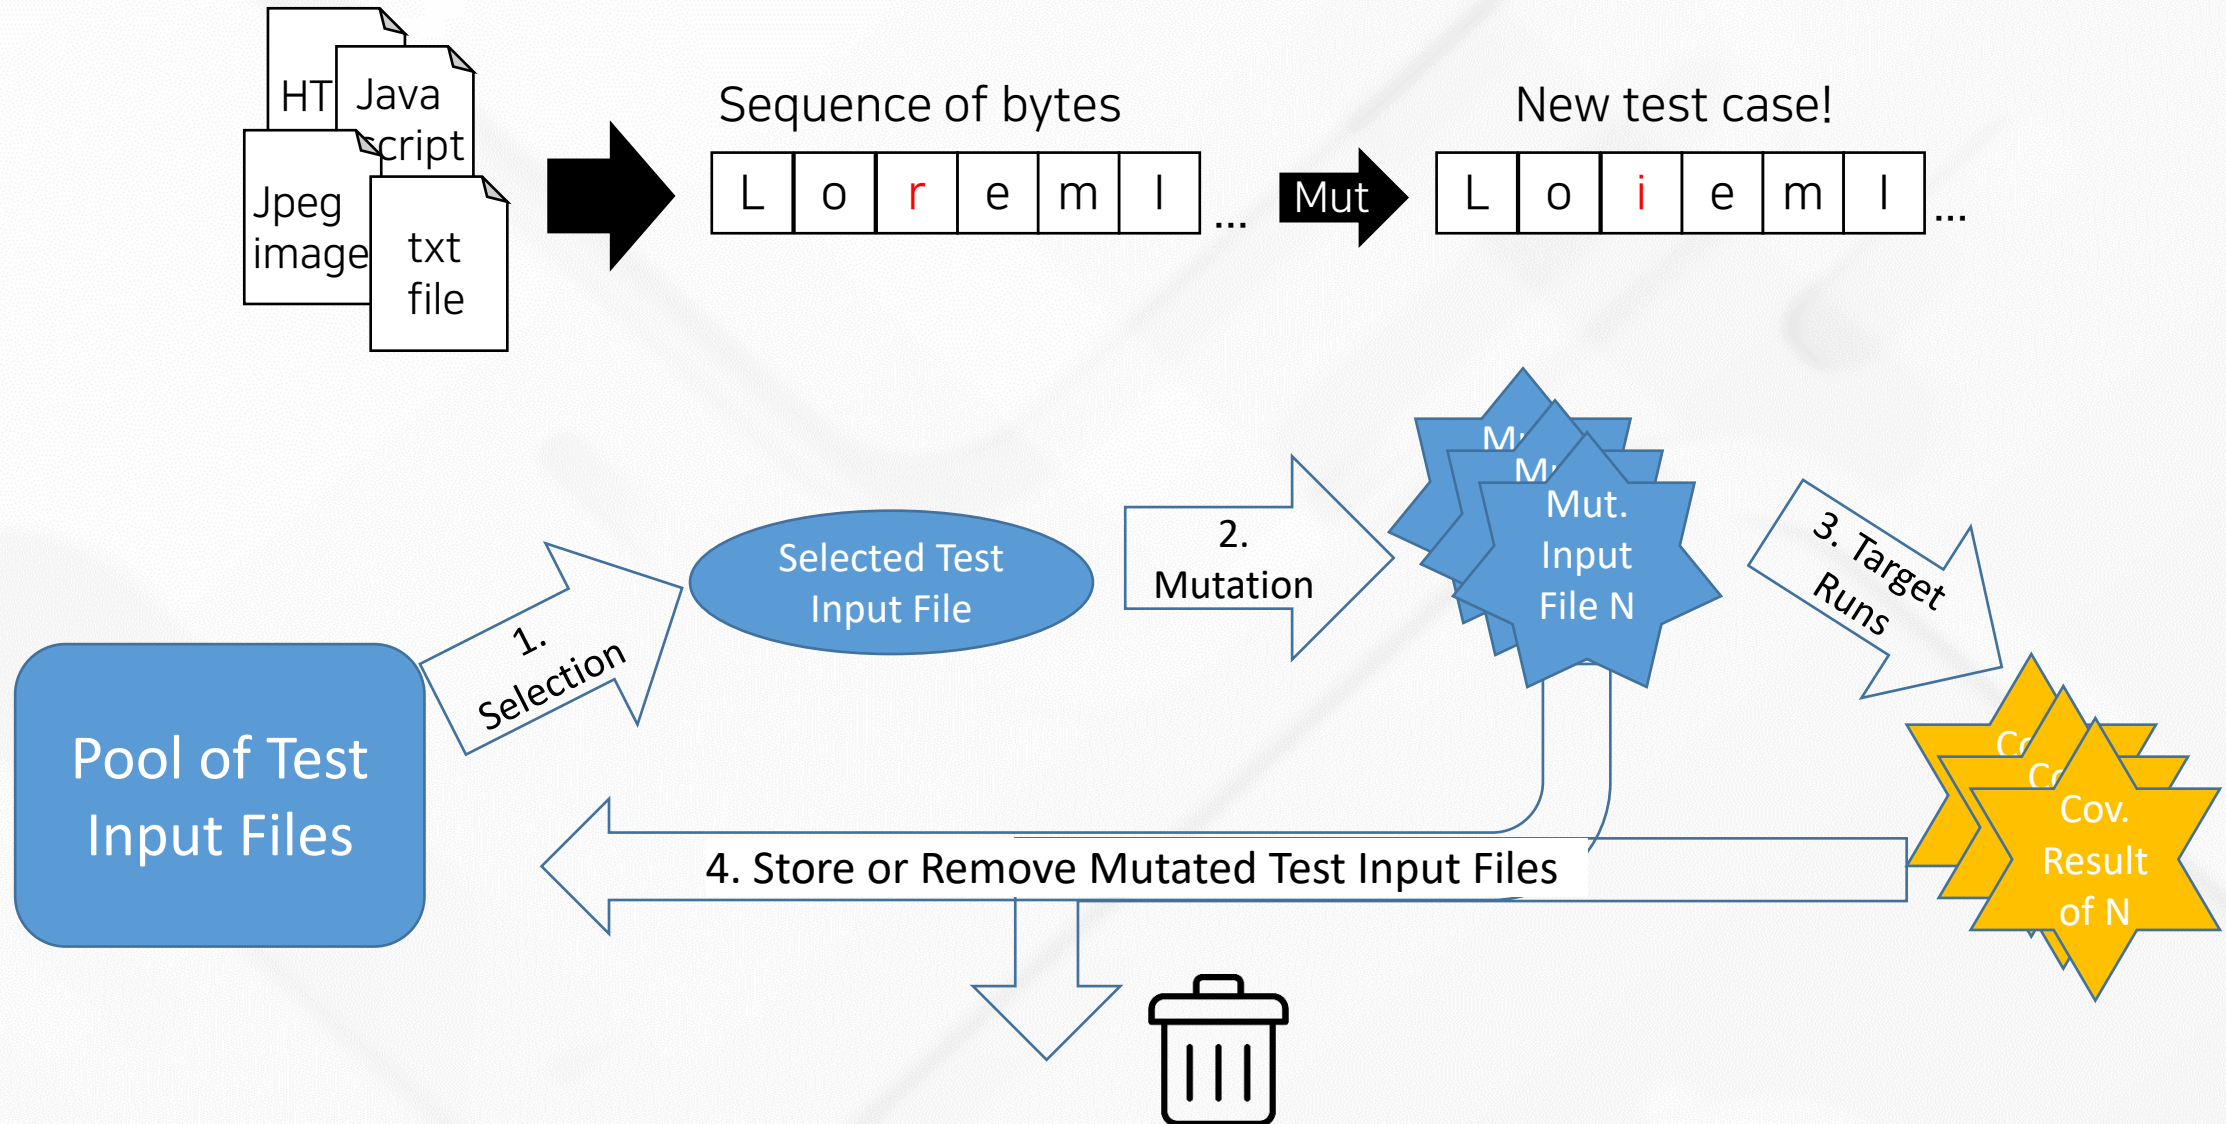
Moonzoo Kim

KAIST

Various SW Testing Techniques w/ Different Cost and Effectiveness

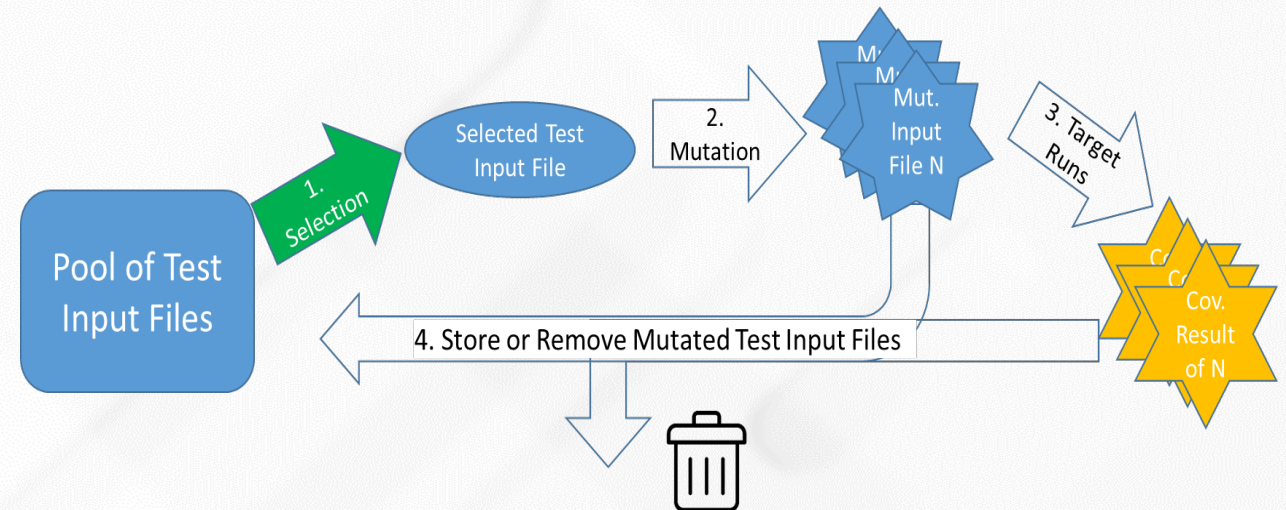


Fuzzing - Automated Test Input Generation via Random Mutation



Fuzzing Challenge #1


Which test input file to select to mutate?



Fuzzer	Heuristic
AFL	Favor test inputs whose execution time and length are short - semantic characteristics of test inputs are ignored)
FairFuzz (ASE 2018), Vuzzer (NDSS 2017)	Favor test inputs such that those test inputs covers hardly covered branches - semantic characteristics of test inputs are a little bit used
CollAFL (SP 2018), Angora (SP 2018)	Favor test inputs that cover many branches whose branch condition statements are executed but these branches are rarely covered - semantic characteristics of test inputs are a little bit used

Fuzzing Challenge #2

Which bytes in the selected test input file should be mutated?



Fuzzer	Heuristic
AFL	Bytes of random # are randomly selected to mutate - No semantic information on the bytes are utilized
FairFuzz (ASE 2018), Profuzzer(SP 2019), GreyOne (USENIX security '20)	Guess which bytes are important based on the runtime information (e.g., if one byte is mutated and the runtime information does not change at all, a fuzzer does not select that byte afterward)
BuzzFuzz(ICSE 2009), Dowser(USENIX security '13), Vuzzer(NDSS 2017), Angora (SP 2018), Matryoska(CCS '19)	Select the bytes which have dependency with "important" branches through Dynamic Taint Analysis (DTA) - Control dependency is not considered

SBST

..... stands for

**Search Based Software
Testing**



Abbreviations.com



**Site map**[Features](#)[Build & Install](#)[Documentation](#)[Tutorials](#)[Papers](#)**Downloads**[Release 4.00c](#)[All releases](#)[Current devel](#)[License](#)**Links**[Repo \(GitHub\)](#)[Donations](#)[Mailing list](#)

AFL++ Overview

AFLplusplus is the daughter of the [American Fuzzy Lop](#) fuzzer by Michał “lcamtuf” Zalewski and was created initially to incorporate all the best features developed in the years for the fuzzers in the AFL family and not merged in AFL cause it is not updated since November 2017.

american fuzzy lop ++2.65d (libpng_harness) [explore] {0}			
process timing		overall results	
run time : 0 days, 0 hrs, 0 min, 43 sec		cycles done : 15	
last new path : 0 days, 0 hrs, 0 min, 1 sec		total paths : 703	
last uniq crash : none seen yet		uniq crashes : 0	
last uniq hang : none seen yet		uniq hangs : 0	
cycle progress		map coverage	
now processing : 261*1 (37.1%)		map density : 5.78% / 13.98%	
paths timed out : 0 (0.00%)		count coverage : 3.30 bits/tuple	
stage progress		findings in depth	
now trying : splice 14		favored paths : 114 (16.22%)	
stage execs : 31/32 (96.88%)		new edges on : 167 (23.76%)	
total execs : 2.55M		total crashes : 0 (0 unique)	
exec speed : 61.2k/sec		total tmouts : 0 (0 unique)	
fuzzing strategy yields		path geometry	
bit flips : n/a, n/a, n/a		levels : 11	
byte flips : n/a, n/a, n/a		pending : 121	
arithmetics : n/a, n/a, n/a		pend fav : 0	
known ints : n/a, n/a, n/a		own finds : 699	
dictionary : n/a, n/a, n/a		imported : n/a	
havoc/splice : 506/1.05M, 193/1.44M		stability : 99.88%	
py/custom : 0/0, 0/0			
trim : 19.25%/53.2k, n/a		[cpu000: 12%]	

The AFL++ fuzzing framework includes the following:

- A fuzzer with many mutators and configurations: afl-fuzz.
- Different source code instrumentation modules: LLVM mode, afl-as, GCC plugin.
- Different binary code instrumentation modules: QEMU mode, Unicorn mode, QBDI mode.
- Utilities for testcase/corpus minimization: afl-tmin, afl-cmin.
- Helper libraries: libtokencap, libdislocator, libcompco.

OSS-Fuzz

 Search OSS-Fuzz[OSS-Fuzz on GitHub](#)

OSS-Fuzz ^

Architecture

Getting started v

Advanced topics v

Further reading v

Bug fixing guidance

Reference v

FAQ v

OSS-Fuzz

[Fuzz testing](#) is a well-known technique for uncovering programming errors in software. Many of these detectable errors, like [buffer overflow](#), can have serious security implications. Google has found [thousands](#) of security vulnerabilities and stability bugs by deploying [guided in-process fuzzing of Chrome components](#), and we now want to share that service with the open source community.

In cooperation with the [Core Infrastructure Initiative](#) and the [OpenSSF](#), OSS-Fuzz aims to make common open source software more secure and stable by combining modern fuzzing techniques with scalable, distributed execution. Projects that do not qualify for OSS-Fuzz (e.g. closed source) can run their own instances of [ClusterFuzz](#) or [ClusterFuzzLite](#).

We support the [libFuzzer](#), [AFL++](#), and [Honggfuzz](#) fuzzing engines in combination with [Sanitizers](#), as well as [ClusterFuzz](#), a distributed fuzzer execution environment and reporting tool.

Currently, OSS-Fuzz supports C/C++, Rust, Go, Python and Java/JVM code. Other languages supported by [LLVM](#) may work too. OSS-Fuzz supports fuzzing x86_64 and i386 builds.

Learn more about fuzzing


This documentation describes how to use OSS-Fuzz service for your open source project. To learn more about fuzzing in general, we recommend reading [libFuzzer tutorial](#) and the other docs in [google/fuzzing](#) repository. These and some other resources are listed on the [useful links](#) page.

Trophies

As of June 2021, OSS-Fuzz has found over [30,000](#) bugs in [500](#) open source projects.

1. How execution paths of test input files are defined and analyzed
2. How to select test input files
3. How to mutate selected test input files

Example of a New Path



- A path is considered **new**, if it covers a new branch, or a branch's hit count is in a new range of hit counts.
 - The list of ranges are as follows:
 - [1, 2, 3, 4-7, 8-15, 16-31, 32-127, 128+]
 - These ranges are called **buckets**.

- For 2 test input files that have the same branch coverage
 - tc1: A-B-C-A-B-D
 - tc2: A-B-C-A-B-C-A-B-D

- Branch hit count of tc1:

Branch ID	b ₁	b ₂	b ₃	b ₄
Hit count	2	1	1	1

- Branch hit count of tc2:

Branch ID	b ₁	b ₂	b ₃	b ₄
Hit count	3	2	2	1

- tc2 covers a **new path** compared to tc1
 - the hit count for b₁, b₂, and b₃ of tc1 and tc2 cover different buckets.

How Path Information is Stored by AFL

- Branch hit count is used to find a new path
- AFL keeps the path coverage data in a 64kB mem. Each byte represent hit count bucket for each branch as follows:

1 hit bucket: 1

2 hits bucket: 1

3 hits bucket: 1

4-7 hits bucket: 0

8-15 hits bucket: 0

16-31 hits bucket: 0

32-127 hits bucket: 0

128+ hits bucket: 0

1 hit bucket: 0

2 hits bucket: 1

3 hits bucket: 0

4-7 hits bucket: 0

8-15 hits bucket: 0

16-31 hits bucket: 1

32-127 hits bucket: 0

128+ hits bucket: 0

...

1 hit bucket: 0

2 hits bucket: 0

3 hits bucket: 0

4-7 hits bucket: 1

8-15 hits bucket: 1

16-31 hits bucket: 0

32-127 hits bucket: 0


128+ hits bucket: 0

1 byte (8 bits)

64K branches

L1 Cache Size


Usually, the size of L1 cache range from 16KB to 64KB. Higher the L1 cache size, Higher is the System Performance in general. Note: In few systems, the size of Instruction Cache is more than the size of Data Cache while the common



- Note: multiple branches may use the same byte
 - the location of a branch's byte is determined by hashing (i.e., hash collision)

Finding the Set of **Favored** Inputs

- Def) a **favored input** for a branch *b*:
 - input that has the **lowest performance score** for *b*:
 - Performance score: $\text{exec_time}(\mu\text{s}) \times \text{length}(\text{byte})$
 - `exec_time`: time it takes for the target program to execute the input in nanoseconds
 - `length`: size of the input in bytes

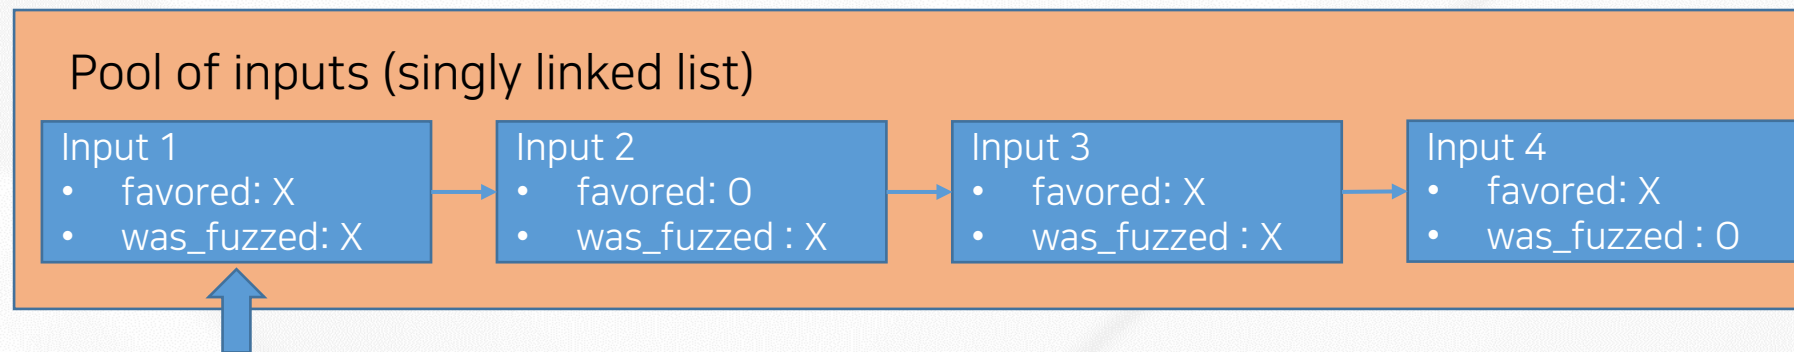



64K favored inputs corresponding to the 64K branches

- AFL selects favored inputs at:
 - the initialization phase after calculating performance score based on the initial seed inputs, and
 - the beginning of each fuzzing cycle

Selecting Input from Pool (1/4)

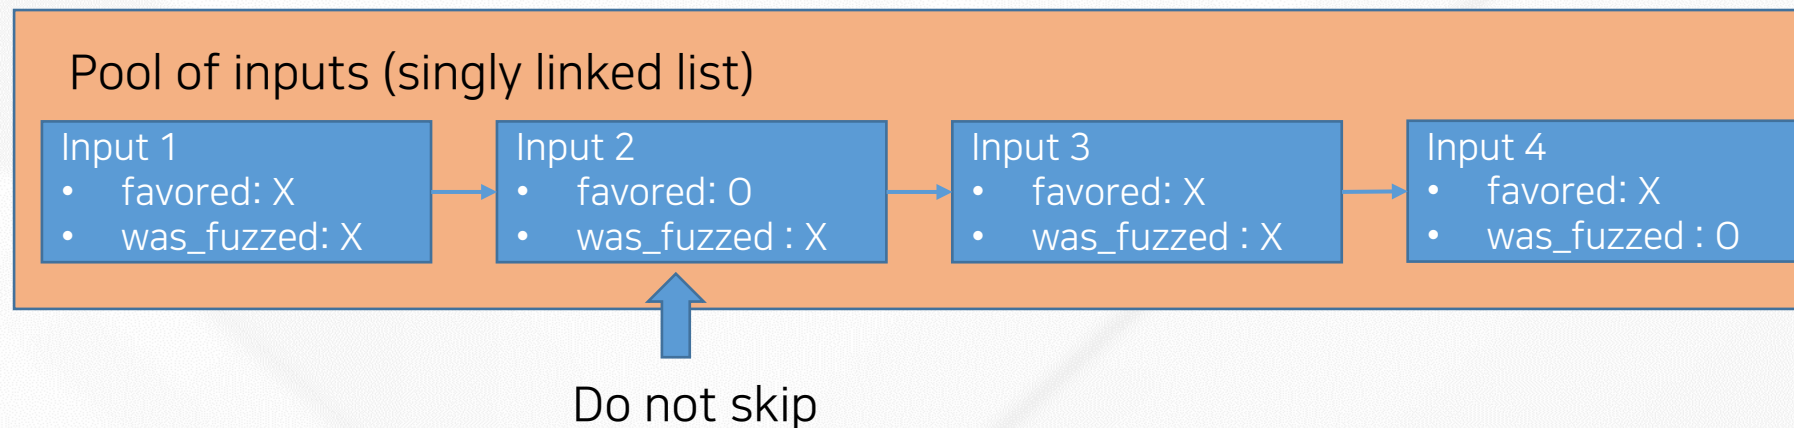
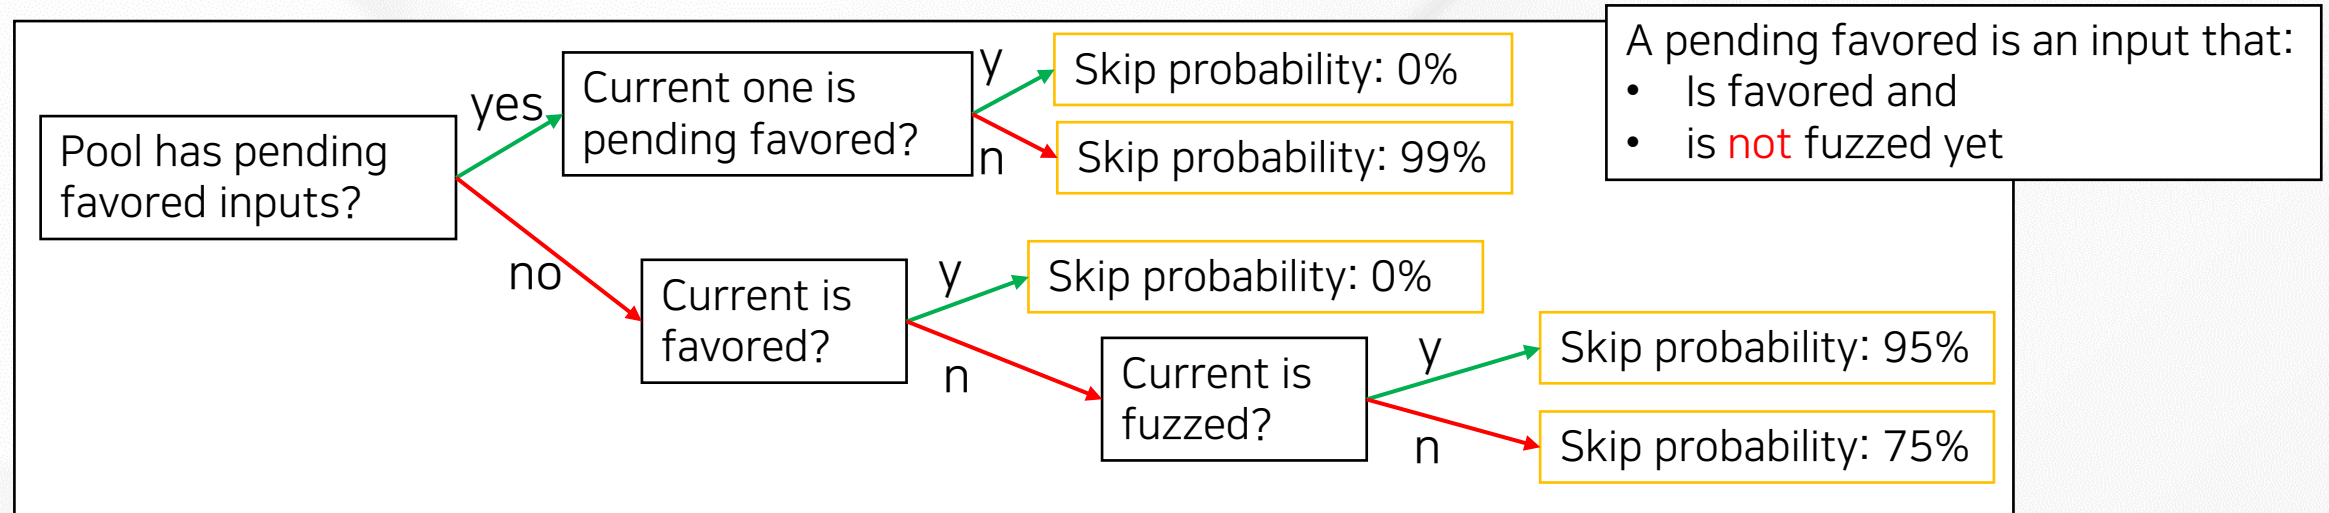
AFL can skip inputs in the pool with probability decided by the following algorithm:



Skip with 99% chance

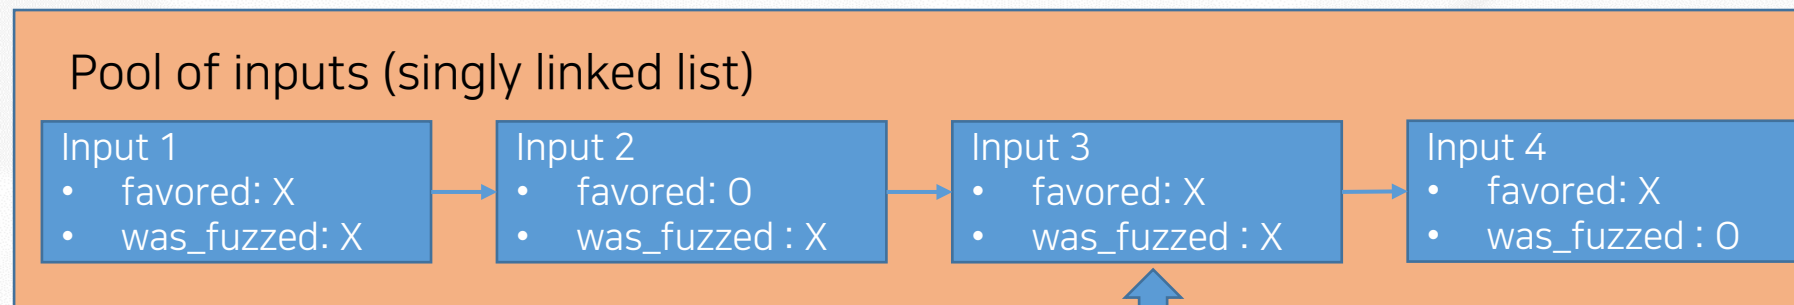
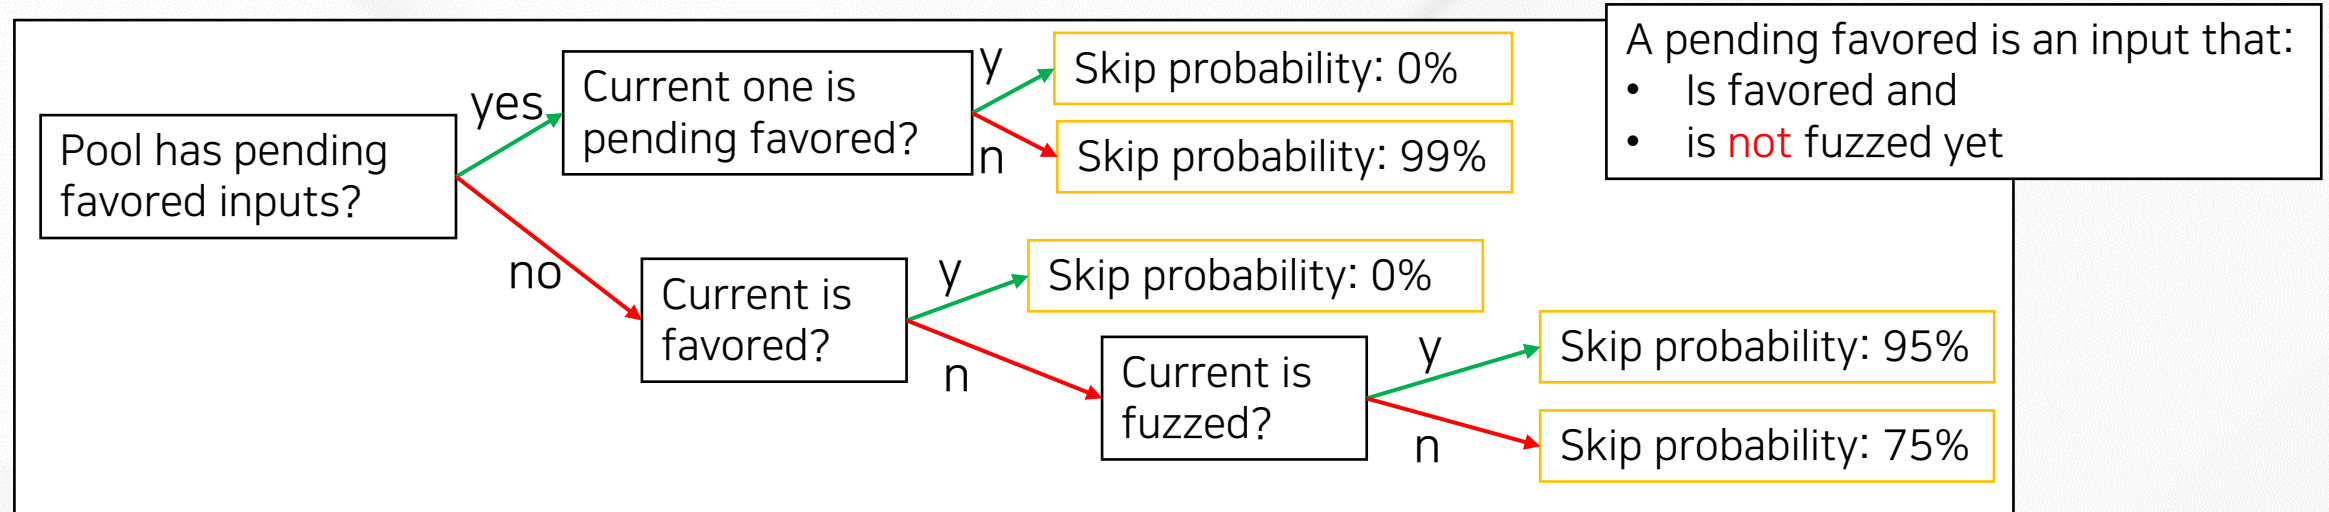
Selecting Input from Pool (2/4)

AFL can skip inputs in the pool with probability decided by the following algorithm:



Selecting Input from Pool (3/4)

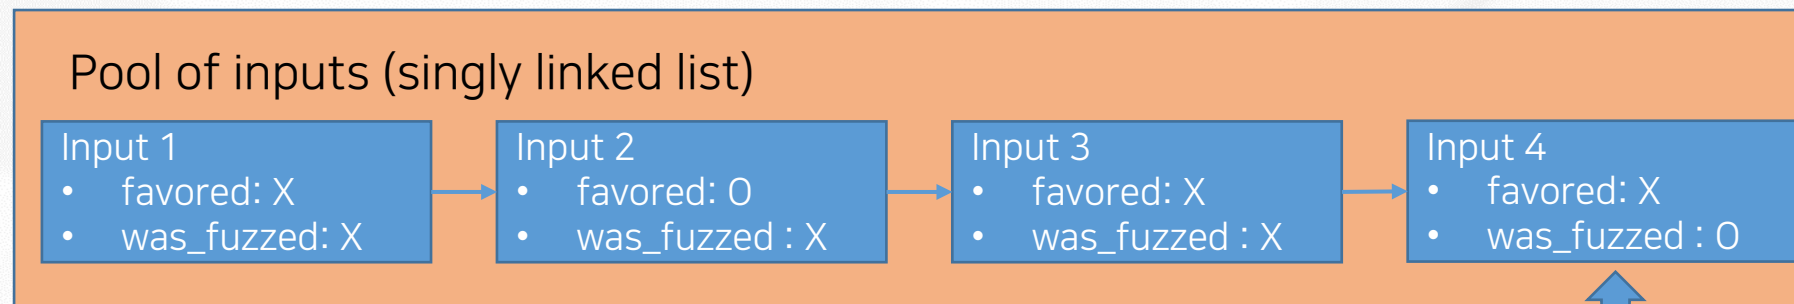

AFL can skip inputs in the pool with probability decided by the following algorithm:



Skip with 75% chance

Selecting Input from Pool (3/4)

AFL can skip inputs in the pool with probability decided by the following algorithm:



Skip with 95% chance

Fuzzing/Mutating Input Bytes

The selected input file is fuzzed using the following fuzzing methods in order:

1. **Bitflip** – flip 1 or 2 or 4 bits of the input
2. **Byteflip** – flip 1 or 2 or 4 bytes of the input
3. **Arithmetic** – add or subtract an integer up to 35 to 8-bit or 16-bit or 32-bit values of the input
4. **Interest** – similar to arithmetic, but overwrite interesting values instead of add or subtract
5. **Extras** – overwrite or insert to the input using user-given or auto-generated terms
6. **Havoc** – makes a random number of modifications to the input using the above 5 methods
7. **Splice** – splice the input with another in the pool and apply havoc

Fuzzing Inputs – bitflip, byteflip

AFL flips L bits at a time, stepping over the input file by S-bit increments. The possible L/S variants are:

- 1/1, 2/1 and 4/1 for bitflip
- 8/8, 16/8, 32/8 for byteflip

Example. The following input (of size 1 byte) is represented in bits as follows:

1	0	1	0	0	0	0	1
---	---	---	---	---	---	---	---

If we apply bitflip 2/1 to the input, it will produce the following fuzzed inputs:

0	1	1	0	0	0	0	1
---	---	---	---	---	---	---	---

1	1	0	0	0	0	0	1
---	---	---	---	---	---	---	---

1	0	0	1	0	0	0	1
---	---	---	---	---	---	---	---

1	0	1	1	1	0	0	1
---	---	---	---	---	---	---	---

1	0	1	0	1	1	0	1
---	---	---	---	---	---	---	---

1	0	1	0	0	1	1	1
---	---	---	---	---	---	---	---


1	0	1	0	0	0	1	0
---	---	---	---	---	---	---	---

Fuzzing Inputs – arithmetic


AFL adds or subtracts integers ranging from 1 to 35 to 8-bit, 16-bit and 32-bit values of the input while stepping over by 8 bits. The possible variants are:

- arith 8/8, arith 16/8, and arith 32/8


Example. The following input (of size 2 bytes) is represented in bits as follows:



If we apply arith 8/8 to the input, it will produce the following fuzzed inputs:




*add 1 to byte 1



*subtract 1 from byte 1



*add 2 to byte 1



*subtract 2 from byte 1

... (up to 35)

... (up to 35)

Fuzzing Inputs – interest (1/2)

AFL overwrites **interesting** values to 8-bit, 16-bit and 32-bit values of the input while stepping over by 8 bits. The possible variants are:


- interest 8/8, interest 16/8, and interest 32/8

The list of interesting values are as follows:


- interesting 8-bit values:
 - -128, -1, 0, 1, 16, 32, 64, 100, 127
- interesting 16-bit values:
 - -32768, -129, 128, 255, 256, 512, 1000, 1024, 4096, 32767
- interesting 32-bit values:
 - -2147483648, -100663046, -32769, 32768, 65535, 65536, 100663045, 2147483647

Fuzzing Inputs – interest (2/2)


Example. The following input (of size 2 bytes) is represented in bits as follows:




If we apply interest 8/8 to the input, it will produce the following fuzzed inputs:




*overwrite -128 to byte 1



*overwrite -128 to byte 2



*overwrite -1 to byte 1



*overwrite -1 to byte 2

Fuzzing Inputs – extras

AFL overwrites or inserts dictionary terms to the input. The dictionary terms can be given by the user or automatically generated by the fuzzer. The possible variants are:

- user extras (over), user extras(insert) - overwrite or insert user given terms
- auto extras (over), auto extras(insert) - overwrite or insert auto generated terms

Example. The following input (of size 8 bytes) is represented in characters:

a	b	c	d	e	f	g	h
---	---	---	---	---	---	---	---

If we overwrite an arbitrary dictionary term “int”, it will produce the following fuzzed inputs:

i	n	t	d	e	f	g	h
---	---	---	---	---	---	---	---

a	i	n	t	e	f	g	h
---	---	---	---	---	---	---	---

a	b	i	n	t	f	g	h
---	---	---	---	---	---	---	---

a	b	c	i	n	t	g	h
---	---	---	---	---	---	---	---

a	b	c	d	i	n	t	h
---	---	---	---	---	---	---	---

a	b	c	d	e	i	n	t
---	---	---	---	---	---	---	---

Fuzzing Inputs – havoc

AFL makes a random number (max 128) of random edits to the input. The number of fuzzed inputs produced is proportional to the performance score of the input.

The list of possible edits are as follows:

- bitflip 1/1
- interest 8/8, 16/8, 32/8
- arith 8/8, 16/8, 32/8
- User extra (over,insert), auto extra (over,insert)
- Set a random byte to a random value
- Remove random number of bytes from random location
- Copy random number of bytes to random location

Fuzzing Inputs – splice

AFL splices together two random inputs from the queue at some arbitrarily selected midpoint and apply havoc

- Applied only after the first full queue cycle ends with no new paths.

